

SAMI HALAWA RIBAS

GCP AI ENGINEER | ENTERPRISE DATA, ML & AGENTIC PLATFORMS

Madrid, Spain · Spanish citizen · Native Spanish · C2 English & Mandarin · Available immediately · linkedin.com/in/samihalawa · github.com/samihalawa

Google Cloud platform ownership

Google Cloud AI engineer and platform architect with 12+ years of software delivery and 5+ years designing data-intensive and AI-enabled systems. **Managed the GCP service estate** across agentic revenue operations, multilingual marketplaces, pricing intelligence, medical AI and regulated-document workflows. I take customer work from discovery and architecture through Python implementation, infrastructure provisioning, deployment, observability and operational handover. The common platform is IAM and VPC foundations; Compute Engine, GKE and Cloud Run runtimes; GCS, AlloyDB AI and BigQuery data; Vertex AI and Gemini Enterprise Agent Platform; and production CI/CD, security and cost control.

BUILD & RUN	DATA & ANALYTICS	AI & AGENTS	OPERATE & SECURE
Compute Engine Cloud Run GKE	GCS · AlloyDB AI BigQuery Pub/Sub · Dataflow	Vertex AI · Gemini Gemini Enterprise Agent Platform · Document AI	IAM · Secret Manager Cloud Build · Artifact Registry Logging · Monitoring · Trace

GCP architecture in production

Founder & GCP AI Systems Architect | AutoClient.ai

Mar 2024 - Present

- **End-to-end workload path.** Ran Python/FastAPI APIs and event-driven workers on Cloud Run, persistent browser and integration workloads on Compute Engine, and GKE-managed worker fleets where Kubernetes networking, scaling and scheduling were required.
- **Enterprise agent platform.** Built multi-agent research, enrichment, qualification, outreach and CRM flows with Vertex AI Gemini and Gemini Enterprise Agent Platform / Vertex AI Agent Platform: ADK-based agents, Agent Runtime, Sessions, Memory Bank, Agent Gateway, tool calling and human approval gates.
- **State and data.** Used AlloyDB AI for PostgreSQL-compatible account state, filtered vector retrieval and agent context; GCS for source artifacts; Pub/Sub, Cloud Tasks and Cloud Scheduler for asynchronous, retrievable handoffs.
- **MLOps and operations.** Owned Cloud Build, Artifact Registry and Cloud Deploy pipelines, IAM/service accounts, Secret Manager, Cloud Logging, Monitoring and Trace across model and application releases.
- **Delivery outcome.** Owned discovery, architecture, coding, deployment, demos and customer iteration; participated in Lanzadera's 2025 startup ecosystem and secured Google Cloud for Startups infrastructure credits.

Founding GCP AI Engineer & Platform Lead | Agents AI Ltd / OULANG

Jul 2025 - Present

- **Internet-scale product delivery.** Built and operate OULANG across web, iOS and Android: 17,262 registered users, 38,857 listings and 89,913 contact reveals as of 3 July 2026.
- **GCP consumer-platform architecture.** Operated Python/FastAPI services and jobs on Cloud Run, persistent workloads on Compute Engine and containerized data-processing fleets on GKE behind Cloud Load Balancing and Cloud CDN; managed GCS media, AlloyDB application state, Pub/Sub events and BigQuery product analytics.
- **Production Vertex AI / Gemini.** Shipped multilingual semantic search at approximately 720 AI queries per day, AI-assisted listing creation and moderation, with grounding, structured outputs and production evaluation.
- **Platform operations.** Owned IAM service accounts and custom roles, VPC controls, Secret Manager, Artifact Registry, deployment automation, Logging and Monitoring; made runtime, scaling and cost decisions across managed and self-managed infrastructure.

Enterprise GCP data and ML architectures

GCP Data & AI Platform Consultant | IWAKY S.L. / AutoPricing

2025 - 2026

- **Pricing-data workload.** Architected the decisive path: marketplace, WhatsApp and ERP inputs into GCS; Pub/Sub and Python/Apache Beam on Dataflow for normalization; AlloyDB for operational state and BigQuery for longitudinal analytics; Vertex AI for matching and pricing signals; Cloud Run for APIs and scheduled jobs.
- **Enterprise delivery.** Converted commercial requirements into architecture, milestones, dashboards and stakeholder-ready outputs for a signed EUR 35,000 scope; EUR 17,500 milestone received and 69 delivery-proof captures produced.
- **Operational ownership.** Automated builds and releases, managed service identities and secrets, and instrumented the ingestion-to-dashboard path for diagnosis and recovery.

GCP ML Engineer & Technical Consultant | Independent / AutoMedical / APOLO

2022 - Present

- **Medical-ML workload.** Managed image datasets and model artifacts in GCS; used GPU compute and Vertex AI Workbench/Pipelines for experimentation and evaluation; registered model versions and deployed endpoints; integrated inference through Python/FastAPI.
- **Inspectible architecture.** Published APOLO-medical-multimodal-instruct as a two-stage image-text system separating visual evidence extraction from downstream reasoning and structured reporting.
- **Consulting delivery.** Owned technical discovery, architecture, rapid prototyping, controlled serving and stakeholder demonstrations for medical AI and data-platform work in English and Spanish.

Regulated document and specialist workflows

GCP Legal AI & Governance Systems | INTLAW / RapidAct / PIME.ai

2025 - Present

- **Regulated-document path.** Document AI and GCS for intake and extraction; Vertex AI Gemini with RAG for cited case analysis; AlloyDB AI for durable state; Cloud Run for review workflows; IAM and KMS for controlled access.
- **Product delivery.** Built omnichannel legal intake, cited case-file analysis and EU AI Act transparency workflows across INTLAW, RapidAct and PIME.ai.

GCP Specialist Workflow & Data Engineer | Umbramed

Jul 2025 - Jul 2026

- **Specialist workflow.** Built a GCS and Document AI ingestion path for medical PDFs, Vertex AI Gemini schema extraction, AlloyDB application state and Cloud Run delivery, with quality reporting in BigQuery.

Selected credentials & role readiness

- **GCP and MLOps.** Machine Learning Operations (MLOps) Specialization - Board Infinity, 2026; Google AI Professional Certificate - Google, 2026; Google Data Analytics Professional Certificate - Google, 2026.
- **End-to-end ML engineering scope.** Designs low-code and coded AI solutions, manages data and models, scales prototypes, serves models, automates ML pipelines and monitors production AI across Vertex AI, BigQuery, GKE, Cloud Run, AlloyDB and Cloud Build.
- **Additional AI engineering.** Agentic AI Engineering; MLOps/LLMOps; Databricks Mosaic AI; LangGraph/LangChain; MCP-powered systems; Microsoft Azure AI Fundamentals (AI-900).

Languages: Spanish (native); English (C2); Mandarin Chinese (C2); Cantonese (conversational). **Location:** Madrid, Spain; remote and periodic European travel. **Start:** Immediate.

Cross-portfolio GCP service-estate ownership

Owned the platform as a working system, not a list of products: identity and least-privilege access, secrets and configuration, artifacts and CI/CD, runtime selection, data lifecycle, logging and monitoring, release recovery, and cost-aware operation across customer and product environments.

OWNERSHIP DOMAIN	GCP SERVICES	ENGINEERING RESPONSIBILITY
Foundation & security	Projects, billing and quotas; IAM, service accounts, VPC, Secret Manager, Cloud KMS	Least-privilege access, configuration, private connectivity, budget and environment ownership
Runtime & serving	Compute Engine, Cloud Run services/jobs/worker pools, GKE Autopilot/Standard, Load Balancing, Cloud CDN	Select runtimes by persistence, scale, networking and GPU needs; operate revisions, health and capacity
Data & events	Cloud Storage, AlloyDB AI, BigQuery, Pub/Sub, Dataflow, Cloud Tasks, Cloud Scheduler	Data lifecycle, operational state, vector retrieval, batch/stream ingestion, analytics and retries
AI, ML & agents	Vertex AI Gemini, Pipelines, Model Registry, Endpoints, RAG Engine, Vector Search; Gemini Enterprise Agent Platform; Document AI	Build, evaluate, deploy and observe models and stateful agents; sessions, memory, tools and document extraction
Delivery & operations	Cloud Build, Artifact Registry, Cloud Deploy, Logging, Monitoring, Trace, Audit Logs, Error Reporting	CI/CD, artifact promotion, release control, SLOs, alerts, tracing, incident diagnosis and auditability

Advanced Python, MLOps and enterprise delivery

- **Python and system design.** FastAPI, Pydantic, AsyncIO, workers, algorithms, REST/gRPC, SQL/PostgreSQL, testing and typed contracts for data-intensive and agentic backends.
- **DevOps and MLOps.** Docker, Kubernetes/GKE, Terraform, Cloud Build, Artifact Registry, Cloud Deploy and Vertex AI Pipelines; evaluation, promotion, rollback, observability and incident response.
- **Enterprise consulting.** Discovery, architecture decisions, delivery plans, workshops, demos, documentation, handover and collaboration with software engineers and data scientists.

Public engineering authority

- **Selected public proof.** 249 public non-fork repositories; VUDA visual UI debugging MCP at 80 stars / 7 forks; 12 Hugging Face technical articles; 111 recorded implementation sessions totalling 125.83 hours across Python, Gemini, agents and deployment.

Engineering and consulting foundation

- Senior Full-Stack Engineer | Damehosting

2019 - 2023

 - **Production engineering.** Built APIs and database-backed platforms; operated Linux services, CI/CD, code review, incident response and developer mentoring.
- AI Instructor & Technical Consultant | Independent

2022 - Present

 - **Technical leadership.** Delivered 125.83 recorded hours across Python, agents, RAG, cloud deployment, observability and AI operations in English and Spanish.
- Founder, Full-Stack Builder & Technical Educator | Megawebs / Megacursos

2009 - 2023

 - **Long-horizon delivery.** Built and operated 70+ web, mobile, automation and analytics projects, establishing the product and customer-delivery base beneath the GCP work.