

# SAMI HALAWA RIBAS

## GCP AI ENGINEER | ENTERPRISE DATA, ML & AGENTIC PLATFORMS

Madrid, Spain · Spanish citizen · Native Spanish · C2 English & Mandarin · Available immediately · [linkedin.com/in/samihalawa](https://linkedin.com/in/samihalawa) · [github.com/samihalawa](https://github.com/samihalawa)

### Google Cloud platform ownership

Google Cloud AI engineer and platform architect with 12+ years of software delivery and 5+ years designing data-intensive and AI-enabled systems. **Managed the GCP service estate** across agentic revenue operations, multilingual marketplaces, pricing intelligence, medical AI and regulated-document workflows. I take customer work from discovery and architecture through Python implementation, infrastructure provisioning, deployment, observability and operational handover. The common platform is IAM and VPC foundations; Compute Engine, GKE and Cloud Run runtimes; GCS, AlloyDB AI and BigQuery data; Vertex AI and Gemini Enterprise Agent Platform; and production CI/CD, security and cost control.

| BUILD & RUN                        | DATA & ANALYTICS                                   | AI & AGENTS                                                             | OPERATE & SECURE                                                                        |
|------------------------------------|----------------------------------------------------|-------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| Compute Engine<br>Cloud Run<br>GKE | GCS · AlloyDB AI<br>BigQuery<br>Pub/Sub · Dataflow | Vertex AI · Gemini<br>Gemini Enterprise<br>Agent Platform · Document AI | IAM · Secret Manager<br>Cloud Build · Artifact Registry<br>Logging · Monitoring · Trace |

### GCP architecture in production

Founder & GCP AI Systems Architect | AutoClient.ai

Mar 2024 - Present

- **End-to-end workload path.** Ran Python/FastAPI APIs and event-driven workers on Cloud Run, persistent browser and integration workloads on Compute Engine, and GKE-managed worker fleets where Kubernetes networking, scaling and scheduling were required.
- **Enterprise agent platform.** Built multi-agent research, enrichment, qualification, outreach and CRM flows with Vertex AI Gemini and Gemini Enterprise Agent Platform / Vertex AI Agent Platform: ADK-based agents, Agent Runtime, Sessions, Memory Bank, Agent Gateway, tool calling and human approval gates.
- **State and data.** Used AlloyDB AI for PostgreSQL-compatible account state, filtered vector retrieval and agent context; GCS for source artifacts; Pub/Sub, Cloud Tasks and Cloud Scheduler for asynchronous, retrievable handoffs.
- **MLOps and operations.** Owned Cloud Build, Artifact Registry and Cloud Deploy pipelines, IAM/service accounts, Secret Manager, Cloud Logging, Monitoring and Trace across model and application releases.
- **Delivery outcome.** Owned discovery, architecture, coding, deployment, demos and customer iteration; participated in Lanzadera's 2025 startup ecosystem and secured Google Cloud for Startups infrastructure credits.

Founding GCP AI Engineer & Platform Lead | Agents AI Ltd / OULANG

Jul 2025 - Present

- **Internet-scale product delivery.** Built and operate OULANG across web, iOS and Android: 17,262 registered users, 38,857 listings and 89,913 contact reveals as of 3 July 2026.
- **GCP consumer-platform architecture.** Operated Python/FastAPI services and jobs on Cloud Run, persistent workloads on Compute Engine and containerized data-processing fleets on GKE behind Cloud Load Balancing and Cloud CDN; managed GCS media, AlloyDB application state, Pub/Sub events and BigQuery product analytics.
- **Production Vertex AI / Gemini.** Shipped multilingual semantic search at approximately 720 AI queries per day, AI-assisted listing creation and moderation, with grounding, structured outputs and production evaluation.
- **Platform operations.** Owned IAM service accounts and custom roles, VPC controls, Secret Manager, Artifact Registry, deployment automation, Logging and Monitoring; made runtime, scaling and cost decisions across managed and self-managed infrastructure.

## Enterprise GCP data and ML architectures

### GCP Data & AI Platform Consultant | IWAKY S.L. / AutoPricing

2025 - 2026

- **Pricing-data workload.** Architected the decisive path: marketplace, WhatsApp and ERP inputs into GCS; Pub/Sub and Python/Apache Beam on Dataflow for normalization; AlloyDB for operational state and BigQuery for longitudinal analytics; Vertex AI for matching and pricing signals; Cloud Run for APIs and scheduled jobs.
- **Enterprise delivery.** Converted commercial requirements into architecture, milestones, dashboards and stakeholder-ready outputs for a signed EUR 35,000 scope, with a EUR 17,500 milestone received.
- **Operational ownership.** Automated builds and releases, managed service identities and secrets, and instrumented the ingestion-to-dashboard path for diagnosis and recovery.

### GCP ML Engineer & Technical Consultant | Independent / AutoMedical / APOLO

2022 - Present

- **Medical-ML workload.** Managed image datasets and model artifacts in GCS; used GPU compute and Vertex AI Workbench/Pipelines for experimentation and evaluation; registered model versions and deployed endpoints; integrated inference through Python/FastAPI.
- **Inspectible architecture.** Published APOLO-medical-multimodal-instruct as a two-stage image-text system separating visual evidence extraction from downstream reasoning and structured reporting.
- **Consulting delivery.** Owned technical discovery, architecture, rapid prototyping, controlled serving and stakeholder demonstrations for medical AI and data-platform work in English and Spanish.

## Regulated document and specialist workflows

### GCP Legal AI & Governance Systems | INTLAW / RapidAct / PIME.ai

2025 - Present

- **Regulated-document path.** Document AI and GCS for intake and extraction; Vertex AI Gemini with RAG for cited case analysis; AlloyDB AI for durable state; Cloud Run for review workflows; IAM and KMS for controlled access.
- **Product delivery.** Built omnichannel legal intake, cited case-file analysis and EU AI Act transparency workflows across INTLAW, RapidAct and PIME.ai.

### GCP Specialist Workflow & Data Engineer | Umbramed

Jul 2025 - Jul 2026

- **Specialist workflow.** Built a GCS and Document AI ingestion path for medical PDFs, Vertex AI Gemini schema extraction, AlloyDB application state and Cloud Run delivery, with quality reporting in BigQuery.

## Selected credentials & role readiness

- **GCP and MLOps.** Machine Learning Operations (MLOps) Specialization - Board Infinity, 2026; Google AI Professional Certificate - Google, 2026; Google Data Analytics Professional Certificate - Google, 2026.
- **Google Cloud Professional Machine Learning Engineer exam-domain experience.** Hands-on architecture maps to all six exam domains: architecting low-code AI solutions; managing data and models across teams; scaling prototypes; serving and scaling models; automating ML pipelines; and monitoring AI solutions across Vertex AI, BigQuery, GKE, Cloud Run, AlloyDB and Cloud Build.
- **Additional AI engineering.** Agentic AI Engineering; MLOps/LLMOps; Databricks Mosaic AI; LangGraph/LangChain; MCP-powered systems; Microsoft Azure AI Fundamentals (AI-900).

**Languages:** Spanish (native); English (C2); Mandarin Chinese (C2); Cantonese (conversational). **Location:** Madrid, Spain; remote and periodic European travel. **Start:** Immediate.

Cross-portfolio GCP service-estate ownership

**Owned the platform as a working system, not a list of products:** identity and least-privilege access, secrets and configuration, artifacts and CI/CD, runtime selection, data lifecycle, logging and monitoring, release recovery, and cost-aware operation across customer and product environments.

| OWNERSHIP DOMAIN      | GCP SERVICES                                                                                                                     | ENGINEERING RESPONSIBILITY                                                                                      |
|-----------------------|----------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| Foundation & security | Projects, billing and quotas; IAM, service accounts, VPC, Secret Manager, Cloud KMS                                              | Least-privilege access, configuration, private connectivity, budget and environment ownership                   |
| Runtime & serving     | Compute Engine, Cloud Run services/jobs/worker pools, GKE Autopilot/Standard, Load Balancing, Cloud CDN                          | Select runtimes by persistence, scale, networking and GPU needs; operate revisions, health and capacity         |
| Data & events         | Cloud Storage, AlloyDB AI, BigQuery, Pub/Sub, Dataflow, Cloud Tasks, Cloud Scheduler                                             | Data lifecycle, operational state, vector retrieval, batch/stream ingestion, analytics and retries              |
| AI, ML & agents       | Vertex AI Gemini, Pipelines, Model Registry, Endpoints, RAG Engine, Vector Search; Gemini Enterprise Agent Platform; Document AI | Build, evaluate, deploy and observe models and stateful agents; sessions, memory, tools and document extraction |
| Delivery & operations | Cloud Build, Artifact Registry, Cloud Deploy, Logging, Monitoring, Trace, Audit Logs, Error Reporting                            | CI/CD, artifact promotion, release control, SLOs, alerts, tracing, incident diagnosis and auditability          |

Advanced Python, MLOps and enterprise delivery

- **Python and system design.** FastAPI, Pydantic, AsyncIO, workers, algorithms, REST/gRPC, SQL/PostgreSQL, testing and typed contracts for data-intensive and agentic backends.
- **DevOps and MLOps.** Docker, Kubernetes/GKE, Terraform, Cloud Build, Artifact Registry, Cloud Deploy and Vertex AI Pipelines; evaluation, promotion, rollback, observability and incident response.
- **Enterprise consulting.** Discovery, architecture decisions, delivery plans, workshops, demos, documentation, handover and collaboration with software engineers and data scientists.

Public engineering authority

- **Open-source engineering and technical publishing.** 249 public non-fork repositories; VUDA visual UI debugging MCP at 80 stars / 7 forks; 12 Hugging Face technical articles; 111 recorded implementation sessions totalling 125.83 hours across Python, Gemini, agents and deployment.

Engineering and consulting foundation

- Senior Full-Stack Engineer | Damehosting

2019 - 2023

  - **Production engineering.** Built APIs and database-backed platforms; operated Linux services, CI/CD, code review, incident response and developer mentoring.
- AI Instructor & Technical Consultant | Independent

2022 - Present

  - **Technical leadership.** Delivered 125.83 recorded hours across Python, agents, RAG, cloud deployment, observability and AI operations in English and Spanish.
- Founder, Full-Stack Builder & Technical Educator | Megawebs / Megacursos

2009 - 2023

  - **Long-horizon delivery.** Built and operated 70+ web, mobile, automation and analytics projects, establishing the product and customer-delivery base beneath the GCP work.